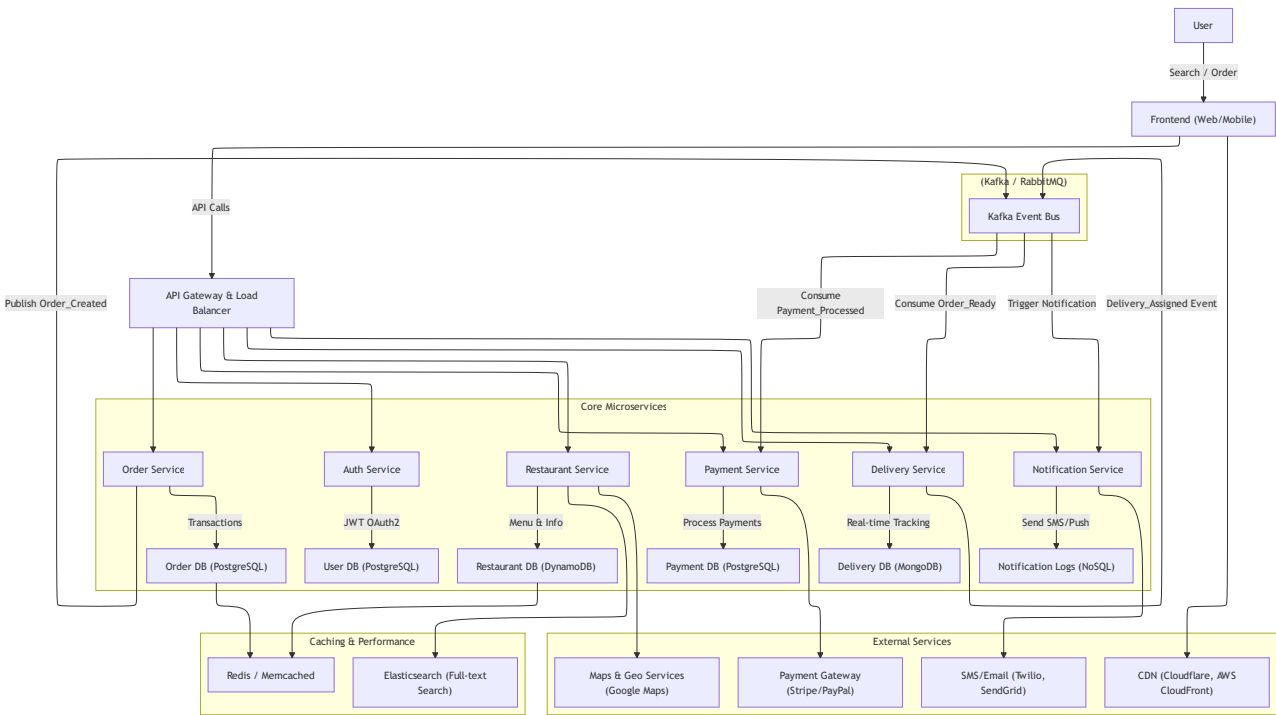


PIZZA DELIVERY SYSTEM (GLOBAL SYSTEM)



1. API Gateway & Load Balancer

- **API Gateway (Kong / AWS API Gateway / Nginx)**
 - Handles authentication, rate limiting, request routing, and caching.
- **Global Load Balancer (AWS ALB / GCP Load Balancer)**
 - Distributes traffic across multiple regional instances.

2. Frontend Clients

- **Mobile App (React Native / Flutter)**
- **Web App (React / Angular)**
- **Admin Dashboard (React / Angular)**
- **Third-Party API Clients (Partner Integration for Orders)**

3. Core Microservices

User & Authentication

- **User Service** (Manages user profiles, authentication, and authorization)
 - **Auth0 / Firebase Auth / AWS Cognito**
 - PostgreSQL for user data storage.
 - JWT / OAuth2 for authentication.

Order Management

- **Order Service** (Handles order creation, updates, and tracking)
 - PostgreSQL (ACID compliance for transactions).
 - Uses Kafka for event-driven updates.

Payment Processing

- **Payment Service** (Handles transactions, refunds, and fraud detection)
 - Integrates **Stripe, PayPal, Razorpay, UPI**.
 - Ensures compliance with **PCI DSS**.
 - Uses **Kafka** for asynchronous order updates.

Restaurant & Menu Management

- **Restaurant Service** (Stores restaurant details and menus)
 - DynamoDB / MongoDB for fast access.
 - Elasticsearch for search functionality.

Kitchen & Order Fulfillment

- **Kitchen Service** (Manages order preparation, estimated time, and workflow)
 - Uses WebSockets for real-time order status updates.

Delivery & Rider Management

- **Delivery Service** (Manages riders and order tracking)
 - Uses **Firebase Firestore** / **MongoDB** for real-time updates.
 - WebSockets / gRPC for real-time communication.

Notification & Customer Engagement

- **Notification Service** (Sends push, SMS, email notifications)
 - Uses **Twilio**, **Firebase Push**, **SendGrid**.

Review & Feedback

- **Review Service** (Manages customer ratings and reviews)
 - Uses PostgreSQL or DynamoDB.

Search & Recommendation

- **Search Service** (Helps users find restaurants and dishes)
 - **Elasticsearch** for full-text search.
 - **Redis** / **Memcached** for caching.

Loyalty & Rewards

- **Loyalty Service** (Handles discounts, coupons, and rewards)
 - PostgreSQL / DynamoDB for storage.

Customer Support

- **Support Service** (Chatbots, helpdesk, and complaint resolution)
 - Uses **Zendesk** / **ChatGPT API**.
-

4. Backend Infrastructure

Database & Storage

- **PostgreSQL** / **MySQL** (For transactional data: orders, users, payments)
- **DynamoDB** / **Cassandra** (For restaurant data, menu caching)
- **MongoDB** / **Firestore** (For real-time delivery tracking)
- **Elasticsearch** (For fast search queries)
- **Redis** / **Memcached** (For caching frequently accessed data)
- **BigQuery** / **Snowflake** (For analytics & reporting)

Event-Driven Architecture

- **Apache Kafka** / **RabbitMQ** (For asynchronous communication)
 - Order events, payment processing, restaurant updates.

Distributed Transaction Management

- **Saga Pattern** (Camunda / Temporal / Orchestration Engine)

Caching & Performance Optimization

- **Redis** / **Memcached** (Reduces DB load)
- **CDN** (Cloudflare / AWS CloudFront) (For static assets)

Rate Limiting & Security

- **API Gateway Throttling**
 - **DDoS Protection** (Cloudflare / AWS Shield)
 - **Data Encryption** (AES-256, TLS 1.2/1.3)
-

5. DevOps & Deployment

Containerization & Orchestration

- **Kubernetes (K8s)** / **Docker** (Manages microservices)
- **Helm Charts** (For Kubernetes deployment)

Continuous Integration & Deployment (CI/CD)

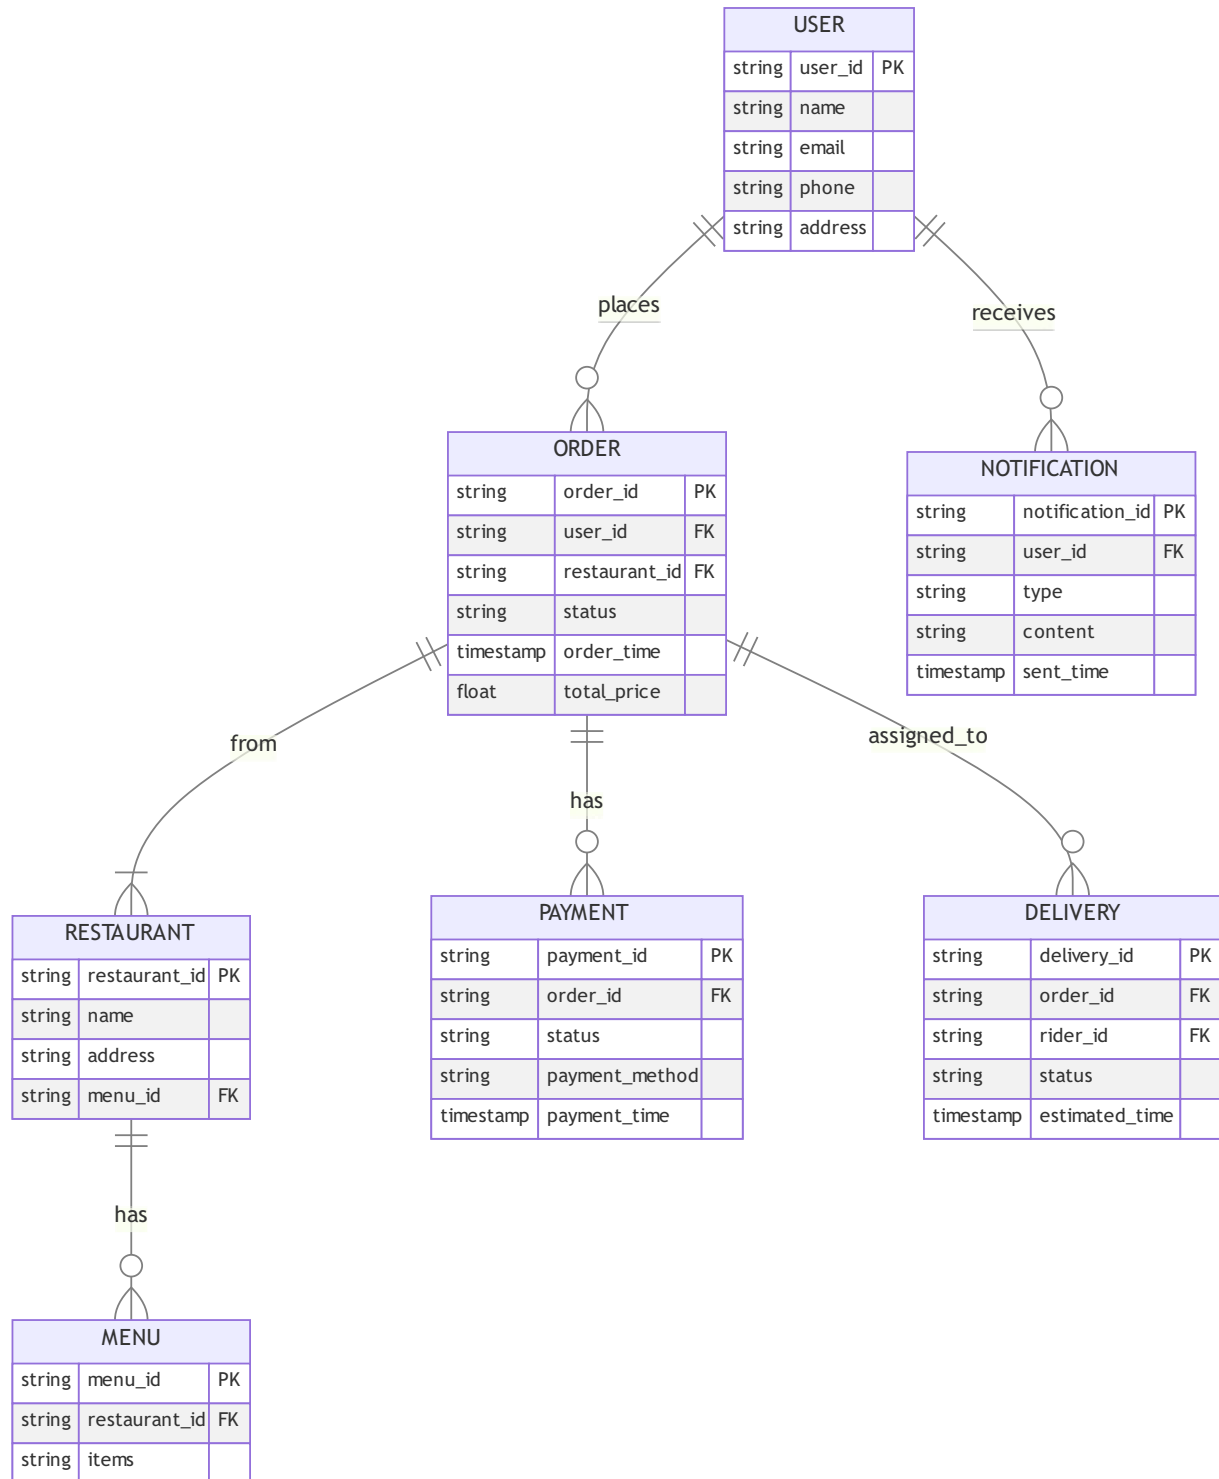
- **GitHub Actions** / **Jenkins** / **ArgoCD** (Automates deployments)

Logging & Monitoring

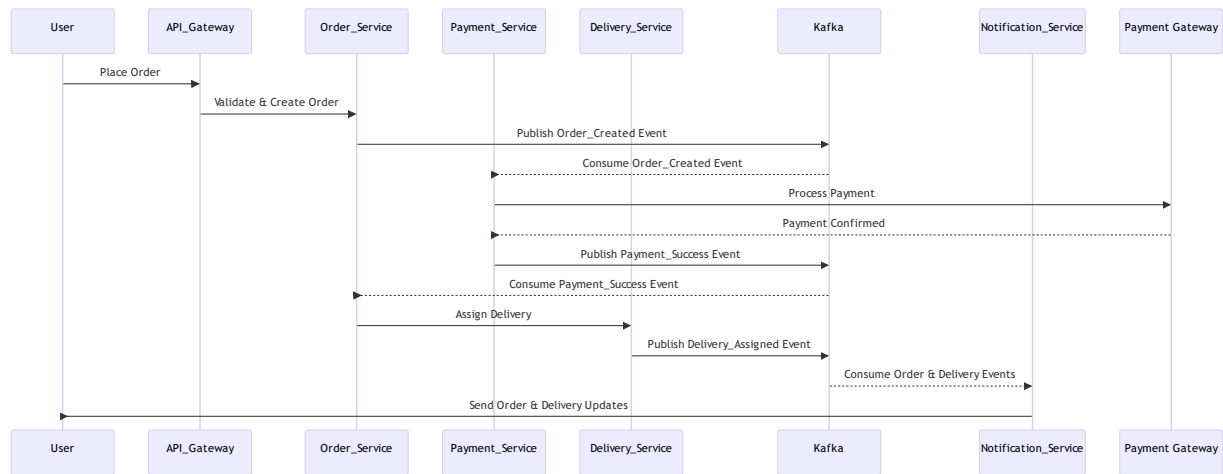
- **Prometheus & Grafana** (Metrics collection & visualization)
- **ELK Stack (Elasticsearch, Logstash, Kibana)** (Log analysis)
- **Jaeger / OpenTelemetry** (Distributed tracing)

6. Third-Party Integrations

- **Payment Gateways:** Stripe, PayPal, Razorpay, UPI.
- **SMS & Email Notifications:** Twilio, SendGrid.
- **Geo Services:** Google Maps API, OpenStreetMap.
- **Analytics & Reporting:** Snowflake, AWS QuickSight.



- ☑ **Frontend Clients** – Apps, Website, Dashboard
- ☑ **Core Microservices** – Orders, Payments, Users, etc.
- ☑ **Infrastructure** – Database, Messaging, API Gateway
- ☑ **Dev Ops & Monitoring** – CI/CD, Logging, Kubernetes
- ☑ **Third-Party Integrations** – Payment, Maps, Notifications



Indexing & Sharding Strategy for the Global Pizza Delivery System

When designing a **global-scale pizza delivery system**, handling **high query loads**, **low-latency search**, and **massive datasets** is crucial. Below, I'll cover **indexing**, **sharding**, **partitioning**, **replication**, and **why specific choices are made**.

1 Indexing Strategy

Indexing helps in **fast lookups** for frequently accessed data. Let's break it down by **service**.

🔗 Order Service (PostgreSQL)

- **Indexes:**
 - **B-Tree Index** on `order_id` (Primary Key) – Faster lookups.
 - **Hash Index** on `user_id` – Efficient filtering by user.
 - **GIN Index** on `status` – Fast lookups for Pending, Completed, etc.
 - **Partial Index** on `created_at` for last 7 days – Optimized queries.
- **Why B-Tree?** PostgreSQL defaults to **B-Tree** because it works well for **range queries** (`BETWEEN`, `<`, `>`).
- **Why Hash Index?** Hash indexes are efficient for **exact match lookups**, such as finding all orders by a `user_id`.

🔗 Restaurant & Menu Service (DynamoDB)

- **Indexes:**
 - **Partition Key** → `restaurant_id`
 - **Global Secondary Index (GSI)** on `city_name`
 - **Local Secondary Index (LSI)** on `cuisine_type`
 - **Elasticsearch Sync** for Full-Text Search on Menus
- **Why GSI & LSI?**
 - **GSI** allows querying restaurants by `city_name`, helpful for geo-based lookups.
 - **LSI** allows efficient filtering by `cuisine_type` within a restaurant.
 - **Elasticsearch** enhances **fuzzy search** for menu items like "Cheese Pizza" ~ "Cheesy Pizza".

🔗 Delivery Service (MongoDB)

- **Indexes:**
 - **2D GeoSpatial Index** on (`latitude`, `longitude`)
 - **Compound Index** on `rider_id`, `status`
 - **TTL Index** on `delivery_time` (Expire after 24 hours)
- **Why GeoSpatial Index?**
 - Helps find **nearest delivery agents** within a given radius.
 - Enables **fast searches** based on **real-time location updates**.
- **Why TTL Index?**
 - Automatically deletes old records to **keep DB size manageable**.

2 Sharding Strategy

Sharding ensures the system can handle **millions of queries per second**. Here's how we'll distribute data.

🔗 Order Service (PostgreSQL)

- **Sharded by user_id (Range-Based Sharding)**
 - Each shard contains **orders of a range of users** (1–10M, 10M–20M).
 - Reduces **hotspotting** since users place orders across different shards.
 - **Replication per region** for **fault tolerance**.
- **Alternative: Hash-Based Sharding**
 - **Pro:** Uniform distribution of load.
 - **Con:** Can't efficiently query "**all orders of a city**".

🔗 Restaurant & Menu Service (DynamoDB)

- **Sharded by restaurant_id (Partition Key)**
 - Each restaurant has its menu data stored separately.
 - **GSI on city_name** ensures we can query "restaurants in a city".
 - **Read replicas in multiple regions** ensure **faster local access**.
- **Why Not Range-Based?**
 - Restaurant names aren't **evenly distributed** (e.g., many start with "Pizza").
 - **Hashing avoids hotspots** in a distributed database.

🔗 Delivery Service (MongoDB)

- **Sharded by city_name (Geo-Sharding)**
 - Each city has its **delivery tracking database**.
 - Within each city, **GeoSpatial indexing** helps find riders.
- **Why Not Hash-Based?**
 - **Location-based lookups** would be inefficient.
 - **GeoPartitioning keeps data close to riders** and improves **real-time updates**.

3 Partitioning & Replication Strategy

Partitioning splits large tables into smaller **manageable units**.

🔗 Order Service (PostgreSQL)

- **Time-Based Partitioning (Monthly)**
 - Orders older than **6 months** move to **cold storage (S3, Glacier)**.
 - Speeds up recent order lookups.
- **Read Replicas per Continent**
 - US, EU, Asia each have a **read-only replica** to **reduce query latency**.

🔗 Restaurant Service (DynamoDB)

- **Partitioned by Region**
 - US-East, US-West, Europe, Asia.
 - **Multi-active replication** ensures **restaurants can update data from any location**.

🔗 Delivery Service (MongoDB)

- **GeoPartitioned & Read-Replicated**
 - India's deliveries stay in India, US in US.
 - **Read-optimized replicas for mobile clients**.

4 Caching Strategy

We need **caching** to avoid **hitting the database too often**.

Service	Cache Strategy	Why?
Order Service	Redis (LRU) for recent orders	Users check recent orders often
Restaurant Menu	Redis / Memcached	Speeds up restaurant & menu queries
Delivery Tracking	In-Memory Cache (10s TTL)	Reduces DB calls for real-time updates
Search & Filters	Elasticsearch	Full-text search for restaurant names

5 Load Balancing & API Gateway

To handle **millions of requests per second**, we use:

Component	Load Balancer / Proxy
API Gateway	Nginx / AWS API Gateway
Backend Services	Kubernetes Ingress + Internal Load Balancers

Caching Layer **Redis Cluster Mode**
Database **Read Replicas & Auto-Scaling**

6 Deployment & Scaling Strategy

Horizontal Scaling

- **Kubernetes (K8s) & Auto-Scaling**
- Each service runs **independently in containers**.
- **Event-driven scaling** → Auto-scale when high load detected.

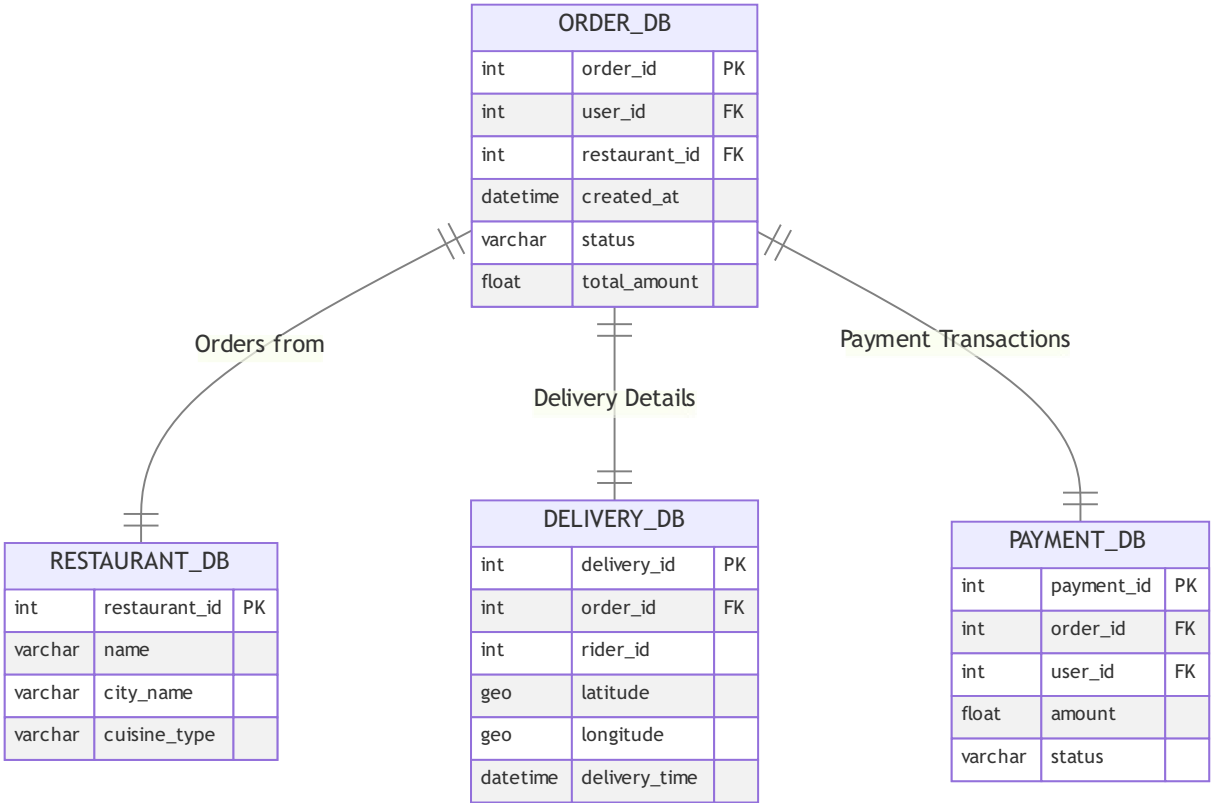
Global Deployment

- **Multi-Region AWS / GCP Setup**
- **Load Balancers** route traffic to nearest server.
- **Edge caching (CDN)** for static assets.

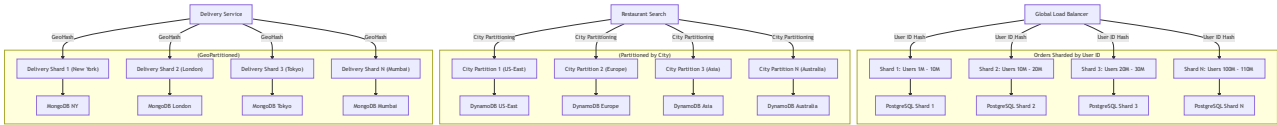
Why This Design?

- ✓ **Optimized Reads** – Using **Indexes, caching, partitioning**.
- ✓ **Low Latency** – **Sharding, Replication, CDN, GeoPartitioning**.
- ✓ **Scalable** – **Event-Driven, Auto-Scaling, Kubernetes**.
- ✓ **Fault-Tolerant** – **Redundant regions, Replication, Kafka**.
- ✓ **Real-time Capable** – **GeoSpatial Indexing, WebSockets, Event Bus**.

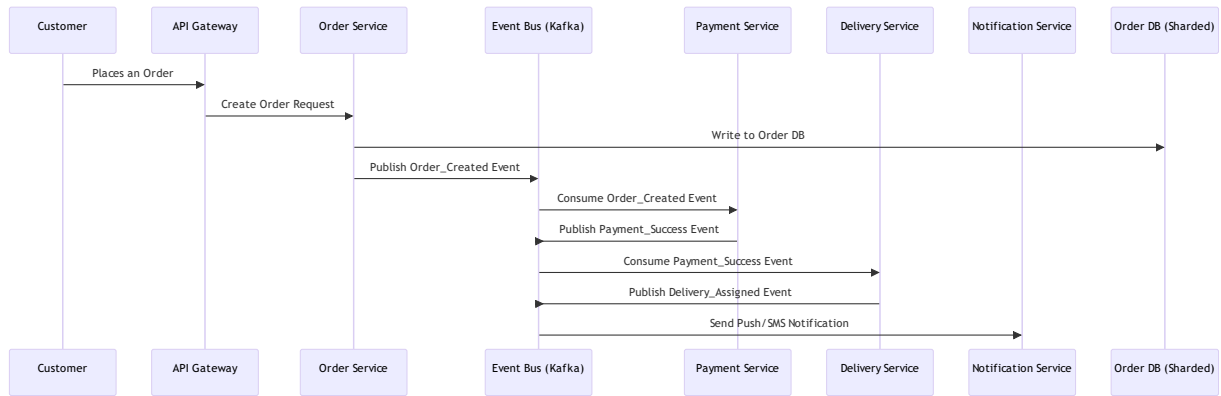
1 Database Indexing Strategy



2 Sharding & Partitioning Strategy



3 Event-Driven Architecture & Replication



Why This Design?

- ✓ **Indexes** optimize search queries across **PostgreSQL, DynamoDB, MongoDB**.
- ✓ **Sharding by User ID, GeoPartitioning, and City-Based Partitions** distribute load.
- ✓ **Event-Driven Architecture with Kafka** reduces synchronous calls & improves scalability.
- ✓ **Read Replicas in Multiple Regions** reduce database load.

Deep Dive into Data Storage & Architecture for a Global Pizza Delivery System 🗺️

A **global-scale pizza delivery system** needs a highly efficient **data architecture** that balances:

- ✓ **Scalability** (handling billions of orders)
- ✓ **Low Latency** (fast search & tracking)
- ✓ **Reliability** (ACID for payments, eventual consistency for tracking)
- ✓ **Cost Efficiency** (hot vs. cold storage)

🔑 1. Key Data Entities & Size Estimations

1 Orders

- **Contains:** Order details, status, payment, delivery info, tracking, etc.
- **Scale:** 500M+ orders/day → ~1TB/day (~365TB/year)
- **Storage Choice:** **PostgreSQL (Sharded) + Cold Storage (S3/GCS)**
- **Partitioning Strategy:**
 - **Sharded by user_id** (even distribution of load)
 - **Partitioned by created_at** (for fast retrieval & archival)
- **Indexing:**
 - **B-Tree index on order_id, user_id** (fast lookup)
 - **GIN index on status (JSONB column)** (query orders by status efficiently)
 - **GeolIndex on delivery_location** (for nearest driver assignment)

2 Users

- **Contains:** Name, email, phone, preferences, addresses, loyalty points
- **Scale:** 500M users (~1.5TB)
- **Storage Choice:** **PostgreSQL + Redis Cache**
- **Indexing:**
 - **Primary Key: user_id**
 - **B-Tree index on phone_number, email**
 - **Cache recent lookups in Redis** (low latency reads)

3 Restaurants

- **Contains:** Name, location, cuisine, ratings, menus, operating hours
- **Scale:** 50M restaurants (~250GB)
- **Storage Choice:** **DynamoDB (NoSQL, GSI on city_name)**
- **Indexing:**
 - **Partition Key: restaurant_id**
 - **Global Secondary Index (GSI) on city_name**
 - **Local Secondary Index (LSI) on cuisine_type**
- **Replication:** Multi-region to keep search fast worldwide

4 Menu

- **Contains:** Menu items, descriptions, prices, availability
- **Scale:** 5B menu items (~2.5TB)
- **Storage Choice:** **MongoDB (NoSQL, Document DB)**

- Indexing:
 - Index on `restaurant_id` (fast lookups by restaurant)
 - Text search indexing for menu search
-

5 Live Order Tracking

- **Contains:** Delivery status, driver location, estimated arrival time
 - **Scale:** 10M+ concurrent deliveries → **360GB/hour**
 - **Storage Choice:** MongoDB (GeoSharded) + Redis for hot lookups
 - Indexing:
 - GeoIndex on `current_location` (fast lookups by area)
 - TTL Index (auto-delete records after 1 hour)
-

6 Payments

- **Contains:** Payment transactions, refunds, fraud detection data
 - **Scale:** 100M+ transactions/day
 - **Storage Choice:** PostgreSQL (ACID Compliance) + Kafka for async processing
 - Indexing:
 - Primary Key: `transaction_id`
 - Index on `user_id`, `order_id`
-

7 Reviews & Ratings

- **Contains:** User reviews, restaurant ratings
 - **Scale:** 1B+ reviews
 - **Storage Choice:** Elasticsearch (Full-text search) + PostgreSQL
 - Indexing:
 - Full-text search index on `review_text`
 - B-Tree index on `restaurant_id`
-

2. Storage Optimization Techniques

1. Sharding Strategy

- PostgreSQL (Orders, Users, Payments)
 - Horizontal sharding by `user_id`
 - Logical partitions on `created_at`
- MongoDB (Live Tracking, Menus)
 - GeoSharding by `latitude`, `longitude`
- DynamoDB (Restaurants, Menus)
 - Partition by `city_name` (ensures load balancing)

2. Caching Strategy

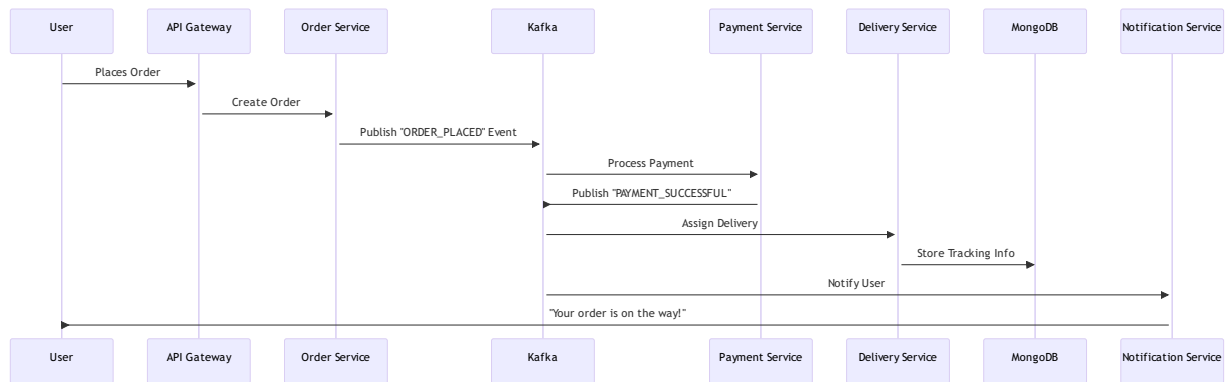
- Redis for user profiles, session management
- CDN (Cloudflare/Akamai) for menu images, static assets
- Elasticsearch for fast menu & restaurant search

3. Archival & Data Retention

- Orders older than 1 year → Cold Storage (S3/GCS)
 - Live tracking data auto-deleted after 1 hour (TTL Index)
-

3. Event-Driven Architecture (Async Processing)

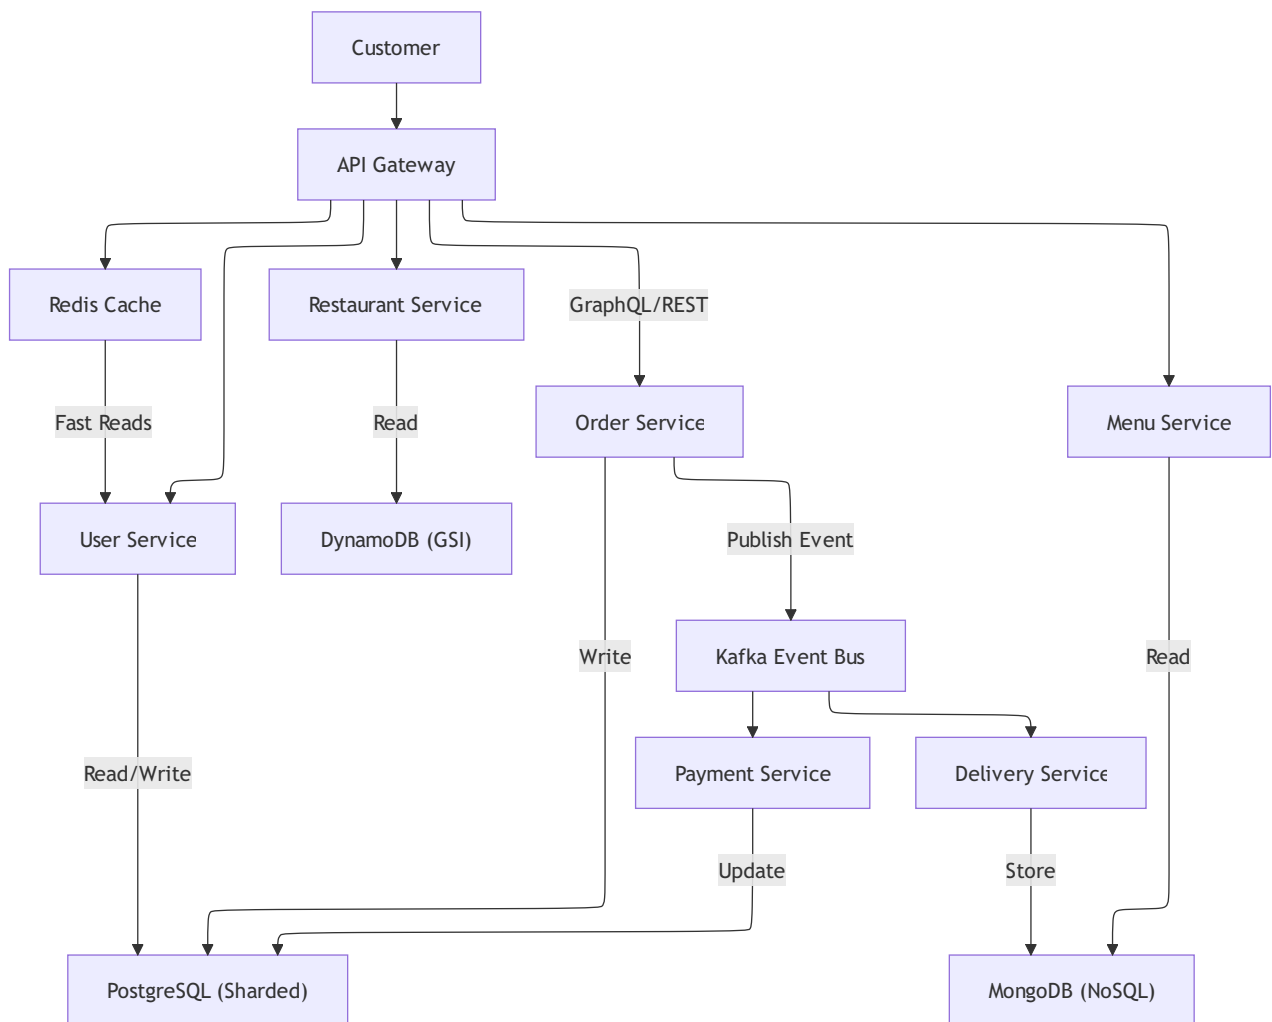
To improve **scalability & performance**, use **event-driven architecture** with Kafka.



☑ Advantages of Kafka-based event processing:

- ✓ **Decoupling Services** (Order, Payment, Delivery work independently)
- ✓ **Scalability** (Millions of orders per second)
- ✓ **Fault Tolerance** (Kafka retains events for replay)

🔗 4. High-Level System Architecture



🔗 5. Deployment & Scaling

🔗 Deployment Model

- Multi-region AWS/GCP Deployment for low latency
- Microservices in Kubernetes (EKS/GKE)
- Auto-scaling with HPA (Horizontal Pod Autoscaler)
- Log aggregation with ELK stack

🔗 Load Balancing

- Global Load Balancer (Cloudflare, AWS ALB)
 - Separate read replicas for PostgreSQL
-

Final Thoughts

Key Takeaways

- ✓ Hybrid SQL + NoSQL for best scalability
- ✓ Sharding & Partitioning to handle billions of records
- ✓ Kafka-based async processing for performance
- ✓ GeoIndexing for live tracking
- ✓ Redis & CDN for ultra-fast access